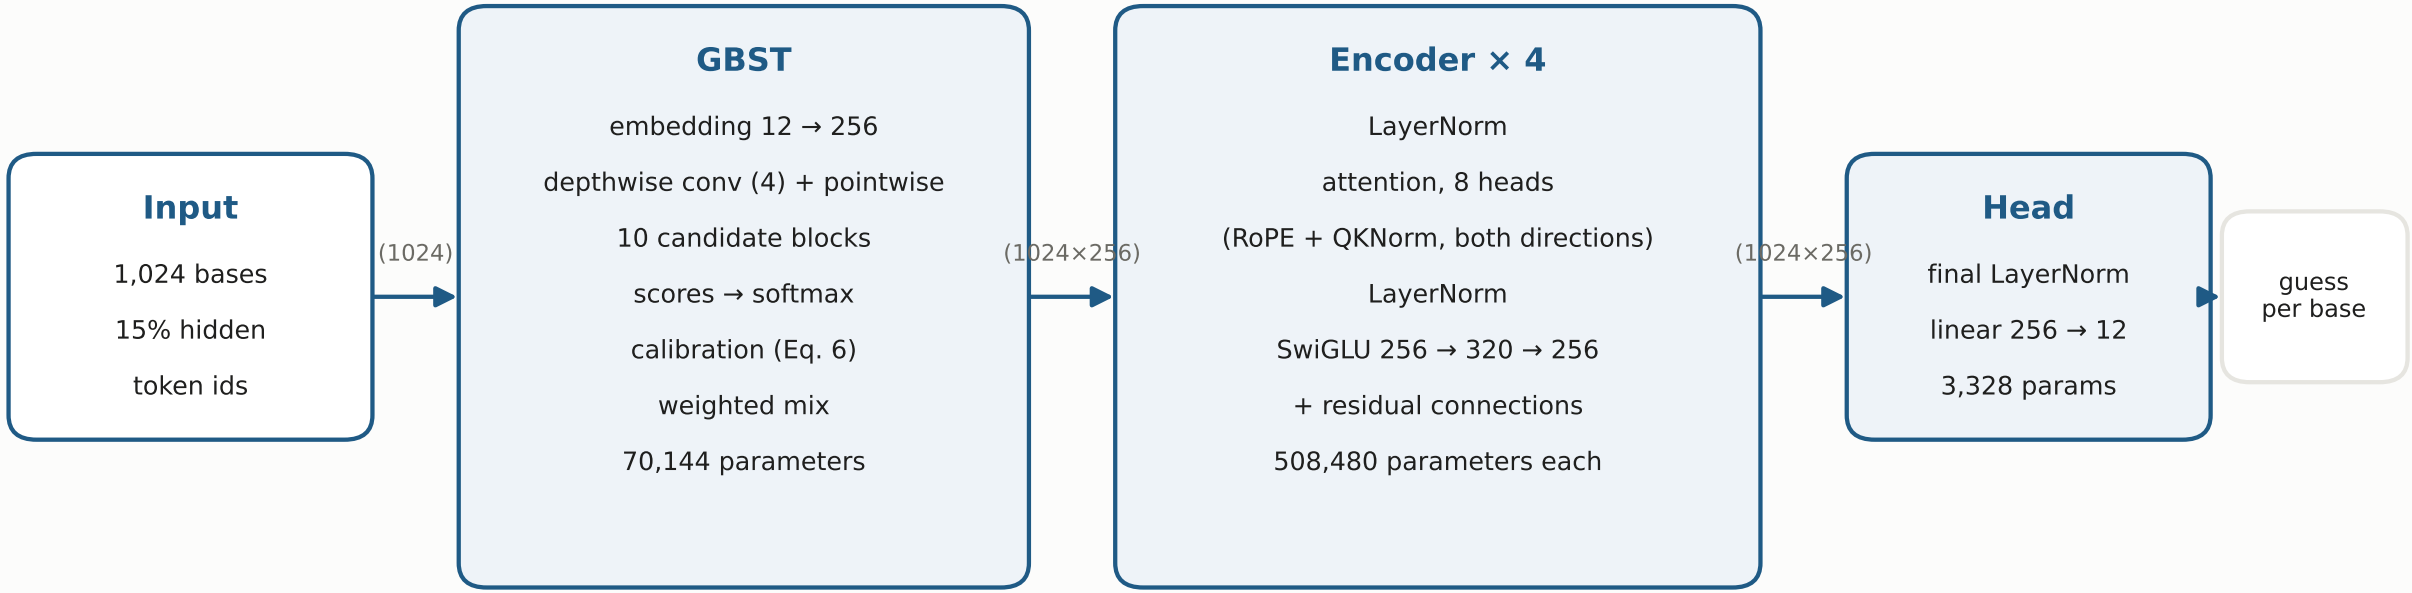
GBST-BERT: 2,107,392 parameters (final model, 10,000 training steps)

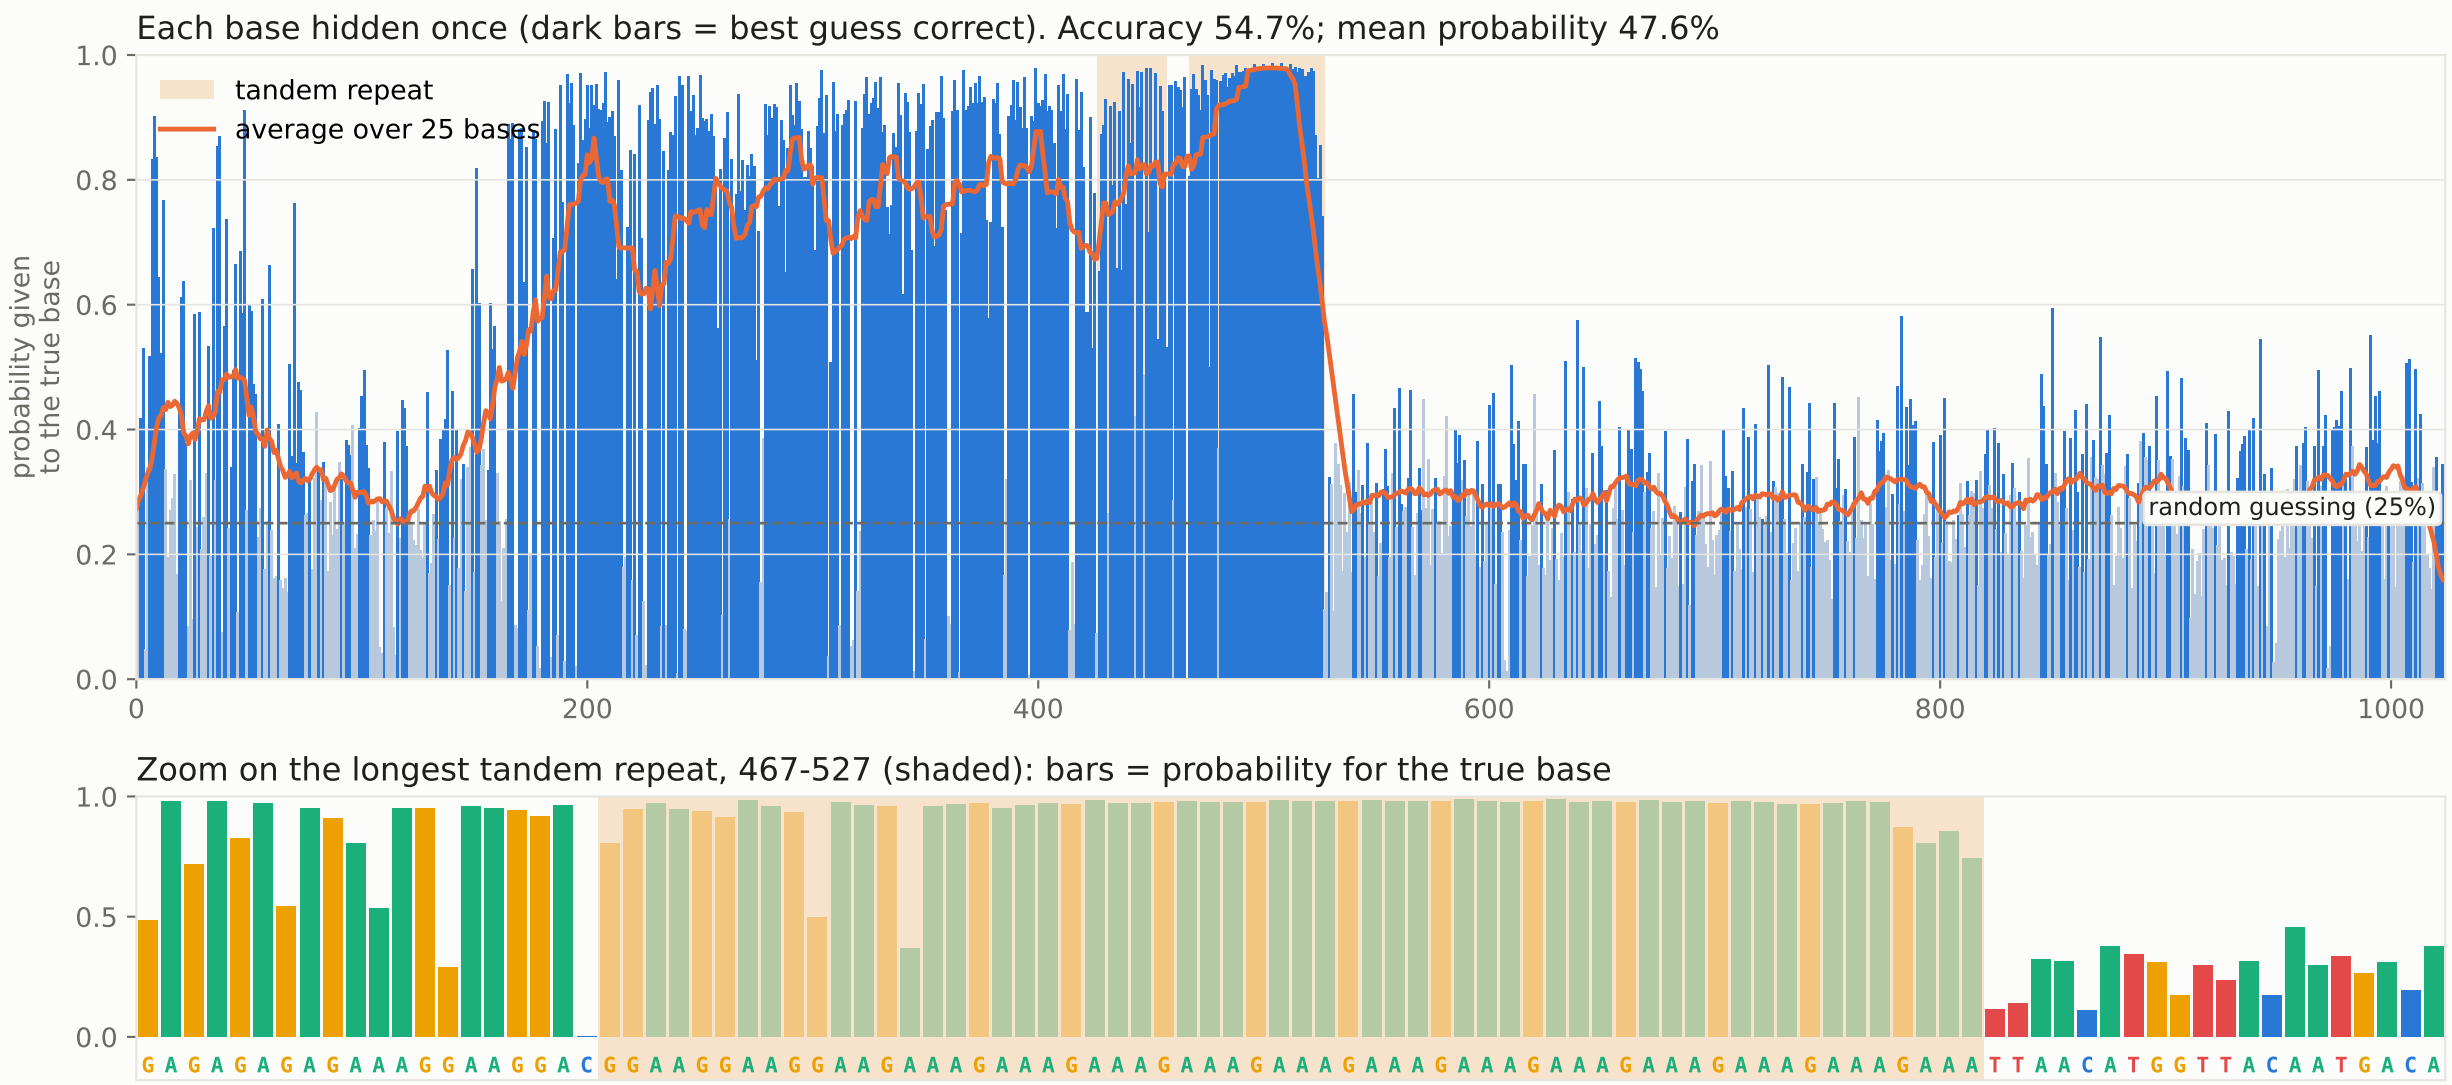


Shapes in brackets: positions × numbers per position. Every base keeps its own position (no downsampling).
Training target: the true base at the hidden positions (loss = cross-entropy).

2 What GBST does along a real DNA sequence



3 How sure the model is, base by base (chr21:33,045,056-33,046,080)



4 Where each base looks: attention per layer (thin lines = the 8 heads)

